

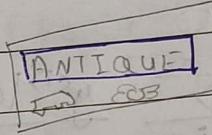
Week 11

Problem Description and Pipeline.
Photo OCR problem.

Photo OCR → Photo optical character Recognition.

* Photo OCR pipeline.

1). Text detection.



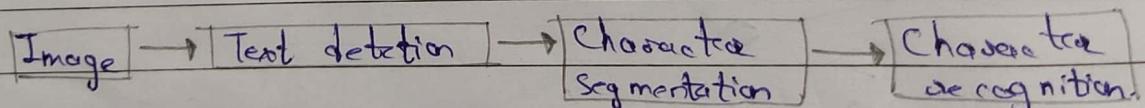
2). Character Segmentation

ANTIQUE MAIL

3). Character classification. $A \rightarrow A$ $N \rightarrow N$ $T \rightarrow T$

In 3rd step it may happen Cleaning, which should be cleaning.
→ We are not considering this 4th step for while.

A system like this is called machine learning pipeline



Machine learning pipeline refers to

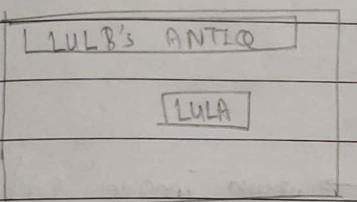
→ A system with many stages/components, several of which may use ml.

--@ Sliding windows —

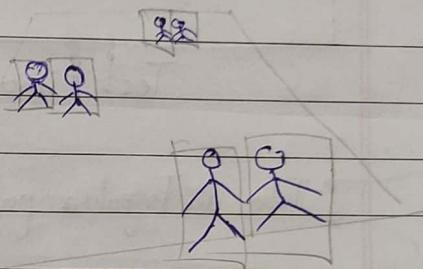
#. How individual component of pipeline works.

. What is sliding windows classifier.

Text detection



Pedestrian detection.



• Aspect ratio is different

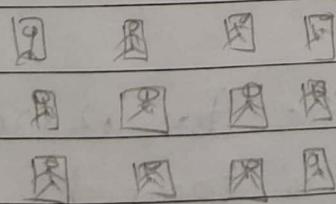
• Aspect ratio is same (approx)

• Height & width of every rectangle may be different.

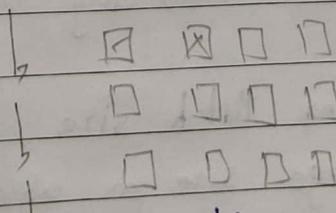
* Supervised learning for pedestrian detection.

x = pixels in 82×36 image patch.

+
+
Do



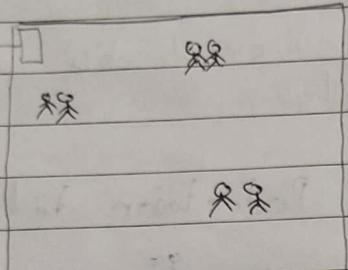
Positive examples ($y=1$)



Negative examples ($y=0$)

Image Patch → contains of pixels.

* Sliding window detection



Test Image.

s1) Take rectangular patch of this image. → sum image path (check for pedestrian)

Slide rectangle & do again. In this way slide through whole test image.

by amount?

It is a parameter →
step-size / stride

(performance is best at 1px (expensive) usually 48 px)

s2) Now increase patch size.

Because previous image would detect pedestrian of particular size (image patch size) only

while
increasing
patch
size →

Take larger Patch size, and resize to 82×36
(patch image size on which model is trained)

This resized image will be passed to classifier
to decide presence of pedestrian in that patch.

s3)

Repeat s1), s2), s3)

So, at end you get all pedestrians spotted.

★ Text detection:

A T R D N D , □ □ □ □ □ □
 B U O W Z W □ □ □ □ □ □

Positive examples ($y=1$)

Negative examples ($y=0$)

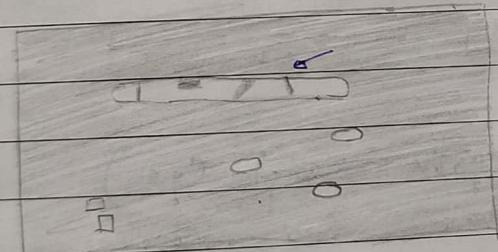
Running sliding window
on lots of little patches like this

LULBA's ANTIQUE MALL

actions
image

Test Image.

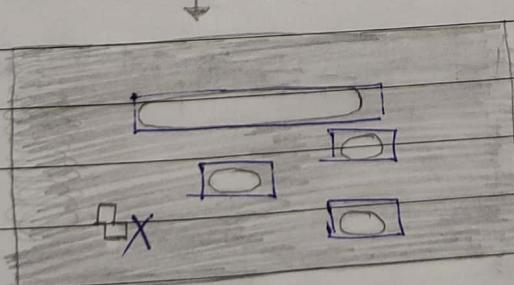
We use sliding window at just one fixed scale (just for purpose of illustration)
(One rectangle size)



• White it classifies has found text.
• Shades of gray (probability of by class)

↓ Expansion operator → Expands white region.

[Colour some neighbouring pixels as white]



• Draw □ around favorable white space. (Aspect ratio must be \square)

Text must be like \square & not like \square (Aspect ratio)

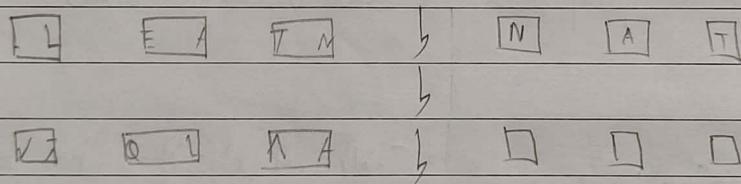
so, we do not consider (\square) such type of white blobs.

We rule out vertical text box.

Q1 Suppose you are running a text detector using $20x20$ image patch. You run the classifier on $200x200$ image. & when using sliding window, you "step" the detector by 4 pixel each time. (For this problem assume you apply the algorithm at only one scale.) About how many times will you end up running your classifier on a single image? (Pick closest one)

- 100
- 400
- about 2,500 times.
- 40,000

Q2) character segmentation • 1D sliding window for character segmentation. look for split b/w two characters.



Positive examples ($y=1$)

These is split

Negative examples ($y=0$)

These is no split.

ANTIQUE MALL

We got this from previous step S/S.

Now perform 1D sliding

titles (only 1 word here)

ANTIQUE MALL

-- (Refer Photo OCR pipeline figure (1) 2 PTO)

Artificial Data Synthesis, (specific problem applicable)

Artificial Data. — — —

Getting lots of Data &

Take low bias algo & big data → Performance ↑.

2.2.2.

ADS

Creating new
Data from
Scratch

We already have
Small training set

Amplify training set
(small ts → large ts)

Character Recognition

A N I T A

* Artificial Data Synthesis for photo OCR:-

(Coloured image is not useful,
It's same as grayscale)

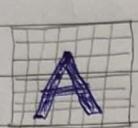
T M B F libus N F s

Real data

Synthetic data.

Using font library, take fonts
& create dataset by giving
some random background.

Synthesizing data by introducing distortion: Text recognition.



Actual dataset



A	A	A	A
A	R	A	R
A	A	A	R
A	A	A	A

By creating distortion.

Synthesizing data by introducing distortion! Speech recognition.

Original audio → Audio on bad cellphone connection.

↳ Noisy background: Crowd.

↳ Noisy background: Machinery.

* Synthesizing data by introducing distortions.

- Distortion introduced should be representation of the type of noise/distortions in the test set.

$$A \rightarrow \begin{matrix} A & A \\ A & A \end{matrix}$$

|| Audio: Background noise, bad cellphone connection.

- Usually does not help to add purely random/meaningless noise to your data.

$$A \rightarrow \begin{matrix} A & A \\ A & A \end{matrix}$$

x_i = intensity (brightness) of pixel i

$$x_i \leftarrow x_i + \text{random noise}$$

Purely random meaningless noise is less likely to be useful.

While adding distortions, think what other meaningful distortions you might add that will cause you to generate additional training examples that are least somewhat representative of the sorts of images you expect to see in your test sets.

^(know org.)

Suppose you are training a LR with m examples by minimizing.

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Suppose you duplicate every example by making two identical copies of it. That is, where you previously had one example $(x^{(i)}, y^{(i)})$, you now have two copies of it, so you now have $2m$ examples, is this likely to help?

→ ✓ No, and in fact you will end up with the same parameters θ as before you duplicate the data.

Discussion on Getting more data.

1) Make sure you have a low bias classifier before expending the effort. (Plot learning curves).

e.g. keep increasing the number of features/ no of hidden units in neural network until you have a low bias classifier.

2) "How much work would it be to get $10x$ as much data as we currently have?"

- Artificial data synthesis

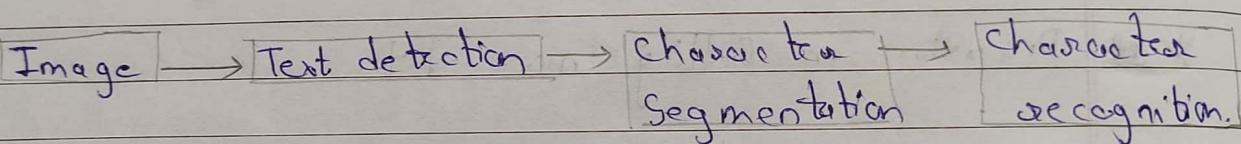
- Collect / label it yourself.

- "Crowd source" (Amazon Mechanical Turk)

-- @ Ceiling Analysis : What part of the Pipeline to work on next.

Ceiling Analysis : Strong advice of what parts of the pipeline might be the best use of your time to work on

* Estimating the errors due to each component
(ceiling analysis)



What part of the pipeline should you spend the most time trying to improve?

Component	Accuracy
overall system	72%
Text detection	89%
char segm	90%
char recgn	100%

In development process, to make decisions on what to do for developing the system is going to very helpful to have single called valued number called evaluation matrix.

Pick character level accuracy:- Given test set image, what is fraction of alphabets or characters in test image that we recognize correctly.

or any single called evaluation matrix.

Image → Text detection → character → characters
↓ ↘
segmentation recognition

- Give correct labels for the text detection part of pipeline manually. So that I get perfect text detection system on my test set.

Run data through rest of pipeline.

(Through char seg. & char. seg.)
Circumfered Text detection 89%.
overall overall

Image → Text detection → characters → character segmentation

✓ both manually supplied.

character segmentation got.

Image → Text detection → character → character
segmentation recognition

✓ Everything manually supplied.

character recognition 100%. (obviously)

1) 2)
17% -

worth
to work
on 2

$$\begin{array}{r} 23 \\ \times 3 \\ \hline 14 \end{array}$$

not worthy
to work.

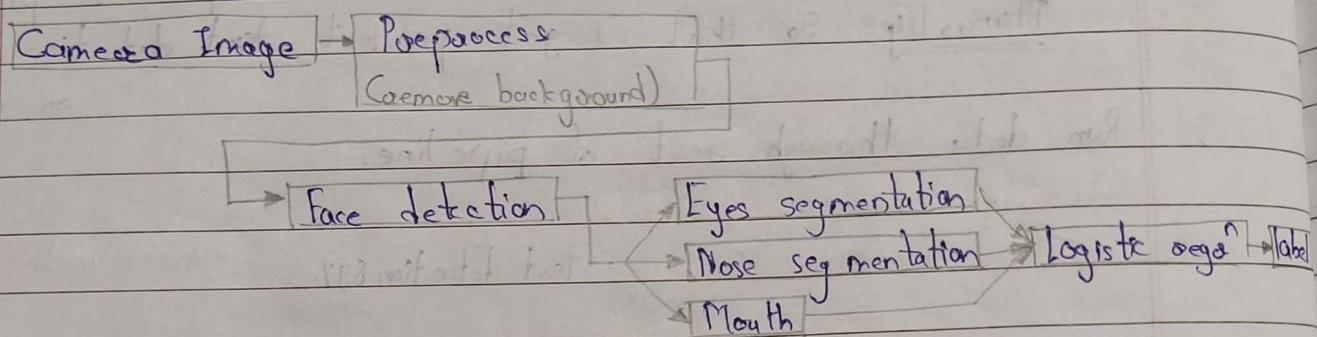
3) 

worthy

We have upside potential
of improving each component

* Another ceiling analysis example.

Face recognition from images (Artificial example)



Component Accuracy.

Overall system 85%.

0.1%

Preprocess 85.1%.

5.9%

Face detection 91%.

4%

Eyes 95%.

1%

Nose 96%.

1%

Mouth 97%.

3%

IR 100%.

Suppose you perform ceiling analysis on a pinched machine learning system, and when we plug in the ground-truth table for one of the components, the performance of the overall system improves very little. This probably means:

- We should dedicate significant effort to collecting more data for that component.

✓ It is probably not worth dedicating engineering resources to improving that component of the system.

✓ If that component is a classifier training using gradient descent, it is probably not worth running gradient descent for 10x as long to see if it converges to better classifier parameters.

• Choosing more features for that component may help (reducing bias), and reducing the number of features for that component (reducing variance) is unlikely to do so.

Instead of going by gut feeling, do ceiling analysis.

-- @ Summary & Thank you.

Summary: Main topics.

Supervised learning: Linear regression, logistic regression, neural networks, SVMs,

Unsupervised learning: K-means, PCA, Anomaly detection.

Special applications/special topics: Recommender system, large scale machine learning.

Advice on building a machine learning system.

Bias/variance, regularization, deciding what to work on next: evaluation of learning algorithm, learning curves, error analysis, ceiling analysis.